

# LOOP UNROLLING

# Instruction-Level Parallelism

➤ Pipelining ?

# Instruction-Level Parallelism

## ➤ Pipelining ?

- improve performance.
- overlapping the execution of instructions.

# Instruction-Level Parallelism

## ➤ Pipelining ?

- improve performance.
- overlapping the execution of instructions.
- This potential overlap among instructions is called **instruction-level parallelism (ILP)** since the instructions can be evaluated in parallel.

# Basic Pipeline Scheduling and Loop Unrolling

- To keep a pipeline full, parallelism among instructions must be exploited by finding sequences of unrelated instructions that can be overlapped in the pipeline.
- To avoid a pipeline stall, a dependent instruction must be separated from the source instruction by a distance in clock cycles equal to the pipeline latency of that source instruction.
- Loop unrolling: replicates the loop body multiple times and adjust the loop termination code.

# Running Example

- This code adds a scalar to an array/vector:

```
for (i=999; i>=0; i=i-1)
    x[i] = x[i] + s;
```

- Assembly Language:

```
Loop: L.D      F0, 0(R1)    ; load the vector element
      ADD.D   F4, F0, F2   ; add the scalar in F2
      S.D     0(R1), F4    ; store the vector element
      DAADUI  R1, R1, #-8  ; decrement the pointer by
                          ; 8 bytes (per DW)
      BNE     R1, R2, Loop ; branch R1 != R2
```

- Register R1: contains the last element of array/vector
- Register R2: is precomputed and 8(R2) contains the first element of the array.

# An Example

Assume following latency all examples

Instruction Producer	Instruction Consumer	Latency
FP ALU op	FP ALU op	3
FP ALU op	Store Double	2
Load Double	FP ALU op	1
Load Double	Store Double	<b>0</b>

- Latency: number of intervening cycles between an instruction that produces a result and instruction that uses the result
- Assume that latency for Integer ops is **zero** and latency for Integer load is **1**

# An Example

<b>Loop:</b>	L.D	F0, 0 (R1)	1
	<b>STALL</b>		2
	ADD.D	F4, F0, F2	3
	<b>STALL</b>		4
	<b>STALL</b>		5
	S.D	0 (R1), F4	6
	DAADUI	R1, R1, #-8	7
	<b>STALL</b>		8
	BNE	R1, R2, <b>Loop</b>	9

This requires **9 Cycles** per iteration



# An Example

## Scheduling

<b>Loop:</b>	L.D	F0, 0 (R1)	1
	DADDUI	R1, R1, #-8	2
	ADD.D	F4, F0, F2	3
	<b>STALL</b>		4
	<b>STALL</b>		5
	S.D	<b>8</b> (R1), F4	6
	BNE	R1, R2 <b>Loop</b>	7

This requires **7 Cycles** per iteration

# An Example

## Unrolling

1	Loop:	L.D	F0, 0(R1)	
2		ADD.D	F4, F0, F2	
3		S.D	0(R1), F4	;drop DSUBUI & BNE
4		L.D	F6, -8(R1)	
5		ADD.D	F8, F6, F2	
6		S.D	-8(R1), F8	;drop DSUBUI & BNE
7		L.D	F10, -16(R1)	
8		ADD.D	F12, F10, F2	
9		S.D	-16(R1), F12	;drop DSUBUI & BNE
10		L.D	F14, -24(R1)	
11		ADD.D	F16, F14, F2	
12		S.D	-24(R1), F16	
13		DADDUI	R1, R1, #-32	
14		BNE	R1, R2, LOOP	

Total Cycles=14+13=27 (L.D ->1 stall, ADD.D-> 2 stall, DADDUI->1 stall).

This requires **6.75 Cycles** per iteration (27/4)

# An Example

Loop :	L.D	F0, 0 (R1)	1
	L.D	F6, - 8 (R1)	2
	L.D	F10, -16 (R1)	3
	L.D	F14, -24 (R1)	4
	ADD.D	F4, F0, F2	5
	ADD.D	F8, F6, F2	6
	ADD.D	F12, F10, F2	7
	ADD.D	F16, F14, F2	8
	S.D	0 (R1), F4	9
	S.D	-8 (R1), F8	10
	DADDUI	R1, R1, #-32	11
	S.D	16 (R1), F12	12
	S.D	8 (R1), F16	13
	BNEZ	R1, R2 LOOP	14

This requires **3.5 Cycles** per iteration. (14/4)

# Loop unrolling (Contd...)

- Suppose the upper bound is  $n$ , and we want  $k$  copies of the loop.
- Don't generate a single unrolled loop!
- Generate a pair of consecutive loops:
  - ❖ First executes the original loop ( $n \bmod k$ )
  - ❖ Second executes the unrolled loop body ( $n/k$ ) times
  - ❖ For large  $n$ , most iterations occur in unrolled loop
- This technique is similar to a technique called **Strip mining**.

# Summary of loop unrolling and scheduling:

- Eliminate the extra test and branch instructions and adjust the loop termination and iteration code.
- Code length increases with additional unrolling. This is issue for embedded processors. This can increase instruction cache miss rates.
- Uses lots of registers. Renaming requires many registers. When registers become scarce: “register pressure”.
- Aggressive unrolling and scheduling and cause a compiler to run out of registers to use for renaming.